

General Regulations.

- Hand in your solutions in groups of two or three people (preferred).
- Practical exercises should be implemented in python and submitted as a jupyter notebook (`.ipynb`). Always provide both the (commented) code as well as the output, and don't forget to explain/interpret the latter.
- If possible, theoretical exercises should be typeset, either by providing your answers in the provided jupyter notebook (use markdown blocks and \LaTeX code for equations) or by submitting a separate `.pdf` file created with \LaTeX . Alternatively, a scan of handwritten notes is also accepted.
- Submit all your files in a single `.zip` archive and upload it using the Physics platform <https://uebungen.physik.uni-heidelberg.de/h/1176>. **Only one person in the group should upload the zip archive:** in the upload interface you will see the option to indicate your exercise-partners.

1 Downsampling an image (9+2pt)

In this exercise we will compare different methods to downscale the image `sign.jpg`. In order to make sure that no interpolation is performed automatically by `matplotlib`, do not forget to set the following option in your jupyter notebook:

```
matplotlib.pyplot.rcParams["image.interpolation"] = None
```

- (a) **Simple subsampling (1pt)** First, load the image in python and plot it at the original resolution. We now want to downsample by a factor of 13 in both directions. As a first approach, start from the pixel at the top-left corner of the original image and then keep only one pixel in every 13 pixels along x and y . Plot and comment your output image.
- (b) **Shifting the image (1pt)** Now repeat, but instead of starting from the pixel at the top-left corner (with coordinates $(0,0)$ in the original image) start from the pixel with coordinates $(4,4)$: is the result different? What if you start from the pixel with coordinates $(8,8)$? Plot and comment your results.

Let us now filter the input image before subsampling it. The size of the filter kernel should depend on the downscaling factor.

- (c) **Box filter (2pt)** We first consider a simple box filter. If the downscaling factor is N , then the box kernel will also have size N (see Fig. 1d and 1e for two examples with $N = 3$ and $N = 5$, respectively). Write a function that returns the box kernel for a given downscaling factor N and then apply the filter to the image `sign.jpg` for $N = 13$. *Hint: use "valid" convolutions to avoid boundary artifacts; similarly to what you have done in the first exercise sheet, use the function `scipy.ndimage.convolve` and the fact that the kernel is separable.*
- (d) **(1pt)** After applying the filter, subsample the image by keeping only one pixel in every 13 pixels starting from the top-left corner. Plot and compare the output with your previous results obtained in point (a). How does the output change if you start subsampling from the pixel with coordinates $(4,4)$ in the filtered image?
- (e) **(1pt)** When we apply the filter and then subsample the image, we are wasting some computations. Why? Have a look at the `torch.nn.functional.conv1d` function in the `pytorch` module: <https://pytorch.org/docs/stable/nn.functional.html>. Which argument of the `conv1d` function would you use to avoid the waste of computations?

- (d) **Lanczos filter (2pt)** The box filter is very fast (because of its small kernel), but it is not a very good low-pass filter and its frequency response can sometimes lead to aliasing when downsampling. The sinc function is the ideal low-pass filter, but it is slower and gives ringing artifacts. Thus, we now consider a commonly used approximation given by the Lanczos windowed-sinc function:

$$\text{Lanczos}_\alpha(x) = \begin{cases} 1 & \text{if } |x| = 0 \\ \frac{\sin(\pi x)}{\pi x} \frac{\sin(\frac{\pi x}{\alpha})}{\frac{\pi x}{\alpha}} & \text{if } 0 < |x| < \alpha \\ 0 & \text{if } |x| \geq \alpha. \end{cases} \quad (1)$$

The two most common choices for the parameter α are 2 and 3, giving a 2-lobed and 3-lobed function, respectively. Write a function computing the normalized Lanczos kernel for a given downscaling factor N and a parameter α . Before to normalize it, its elements should be given by:

$$\text{Lanczos}_\alpha\left(\frac{i}{N}\right) \quad \text{with } i \in \{-N\alpha, \dots, N\alpha\} \quad (2)$$

In Fig. 1f, we show an example of a normalized kernel for $\alpha = 2$ and $N = 3$. More examples and details can be found on pages 10-11 of <http://www.realitypixels.com/turk/computergraphics/ResamplingFilters.pdf>. Using your implemented function, print the Lanczos kernel for $N = 9$ and $\alpha = 2$.

- (e) **(1pt)** Apply the Lanczos filter to the image `sign.jpg` for $N = 13$ and $\alpha = 3$ and then subsample the image. Plot and comment your results.
- (e) **High-frequency image (2pt - Bonus)** Downscale the image `high-frequency.png` by a factor of $N = 5$ with the three methods you implemented previously: sub-sampling using the box filter, the Lanczos filter and no filter at all. Save the three outputs as images and comment why the results are different. *Note that we ask you to save your outputs because if you plot your results with the matplotlib option “`image.interpolation`” set to `None`, matplotlib will likely perform a basic subsampling during the visualization and this could result in plots with aliasing artifacts because of the high frequency content of this image.*

